

Built to Benefit Everyone: Our Plan

Auteur: Sam Altman (CEO) & Jakub Pachocki (Chief Scientist), OpenAI ? Datum: 8 juni 2026

Originele bron: <https://openai.com/nl-NL/index/built-to-benefit-everyone-our-plan/>

Since the inception of OpenAI, our charter has held a singular commitment: to ensure that artificial general intelligence benefits all of humanity. Today, we are outlining the third phase of our roadmap. We are transitioning from proof-of-concept chatbots to ubiquitous, transformative cognitive infrastructure.

Throughout history, foundational general-purpose technologies—from the steam engine to electricity and the internet—fundamentally restructured human civilization. Electricity did not merely light up factories; it empowered every household, transformed medicine, and raised the baseline standard of living for billions of people. Artificial general intelligence must follow this exact path of universal distribution.

Three Strategic Pillars for Phase Three

To realize our mission, our near-term engineering and operational strategy is centered around three pillars:

1. The Automated AI Researcher

The first milestone of Phase Three is deploying an automated AI scientist capable of performing independent, peer-reviewed scientific discovery. By utilizing advanced reasoning models and autonomous experimental design, these systems will accelerate progress in materials science, fusion energy, room-temperature superconductors, and genomic medicine. This will compound our ability to solve humanity's most stubborn challenges.

2. Broad Economic Growth with Distributed Gains

Technology must not become an engine of extreme capital concentration. We are actively structuring our commercial deployment and governance frameworks so that the wealth generated by AGI is distributed globally. This includes innovative dividend mechanisms, public compute grants for developing nations, and universal access tiers that guarantee advanced cognitive tools to everyone regardless of income.

3. A Personal AGI for Every Person on Earth

Our ultimate product vision is not a centralized corporate oracle, but a deeply personal, highly aligned AGI assistant for every individual. Every person will have access to an expert tutor, a creative collaborator, an advocate, and an assistant that works tirelessly to expand their personal capabilities and economic agency.

Avoiding Centralized Monopolies

We reject the premise that superintelligence should be locked inside a closed fortress controlled by a handful of corporate executives or state planners. A singular, hyper-centralized superintelligence would represent an unacceptable concentration of power. By democratizing access through open APIs, modular local runtimes, and distributed computing, we ensure that power remains in the hands of individuals, small businesses, and open societies.

International Governance and Coordinated Safety

At the same time, democratization cannot mean lawlessness. The capability thresholds required to train frontier models require colossal computational infrastructure-gigawatts of energy and millions of specialized chips. This physical reality makes frontier AI uniquely observable and governable.

We reiterate our call for an International Frontier AI Governance Framework:

- (o) Independent safety verification before models exceeding critical compute thresholds are trained or released.
- (o) Coordinated pause protocols: if a frontier system demonstrates dangerous autonomous replication or cyber-offensive capabilities, all participating labs must agree to pause and share diagnostic data.
- (o) Global compute accounting to ensure that high-risk training runs cannot be concealed in unregulated jurisdictions.

The Road Ahead

The transition to AGI will be the most consequential transition in human history. It will bring challenges, disruptions, and difficult ethical choices. But if we act with transparency, broad cooperation, and unwavering dedication to the public good, artificial intelligence will inaugurate an unprecedented era of human flourishing.